

Deep Learning for Computer Vision

Assignment Sheet 3 - Due 24.12.17 23:59

Quick exercise Q&A: 19.12.17, 14:30

Lecture evaluation: 15.12.17, 11:20

Next exercise group: 9.1.18, 13:00

Exercise 3.1 Derivatives (20 points)

This exercise closes the (small) gap in the lecture in the backpropagation algorithm. Let $f : \mathbb{R}^N \rightarrow \mathbb{R}^M$ and $g : \mathbb{R}^M \times \mathbb{R}^M \rightarrow \mathbb{R}$. Compute the derivative of $h : \mathbb{R}^N \rightarrow \mathbb{R}$, $h(\mathbf{x}) := g(f(\mathbf{x}), f(\mathbf{x}))$ in terms of the derivatives of f and g .

Exercise 3.2 Backpropagation (10 + 15 points)

The goal of this exercise is to see and understand how to do backpropagation on a minimalistic example of two-layered fully connected network. Note, in the illustration below, the input is on the bottom, output on the top. Pairs of numbers show (layer index, neuron index). Indices for w s show input neuron on bottom, connected output neuron on top.




Figure 1: A two layer regression fully-connected architecture

The network calculates the function $f : \mathbb{R}^3 \rightarrow \mathbb{R}$ given by

$$f(\mathbf{x}; \mathbf{w}) = \sum_{i=0}^2 w_{1,i}^{2,0} \left(\sum_{j=0}^2 w_{0,j}^{1,i} x_j \right) \quad \text{and loss} \quad E(\mathbf{w}) = \frac{1}{L} \sum_{l=1}^L \|d_l - f(\mathbf{x}_l; \mathbf{w})\|_2^2. \quad (1)$$

- Calculate the partial derivative of the loss with respect to the weight $w_{1,1}^{2,0}$ (depicted in red).
- Calculate the partial derivative of the loss with respect to the weight $w_{0,1}^{1,0}$ (depicted in blue).

Exercise 3.3: Activation functions and batch normalization (10 + 10 + 10 + 10 points)

The purpose of this exercise is to show how one can try different activation functions (PReLU) and batch normalization to get better results or faster convergence.

- Write your own Tensorflow leaky ReLU layer with a parameter $\alpha \geq 0$. Reminder, the leaky ReLU is defined as

$$\text{LeakyReLU}_{\alpha}(x) = \begin{cases} x & x > 0 \\ \alpha x & \text{otherwise} \end{cases} \quad (2)$$

- Train the architecture from Exercise 2.2, but replace the activations with leaky ReLU units (choose the same α for all layers for simplicity). Play around with different values of α and draw graphs with test and train loss/accuracy.
- Now, switch to parametric ReLU, i.e. train α as if it were another parameter. What's the optimal value, and how does it compare to the previous experiments?
- Add batch normalization (as described in the lecture) before the activations. How does it influence training and testing loss? How does it influence convergence?
- (Bonus) Add (non-trivial) convolutional layers to your network, and see how many you can stack while still being able to train the network. Use correct initializations, try e.g. weight regularization/dropout to prevent overfitting. Bonus points: $\min(20, \text{number of working layers divided by } 2)$. You might want to switch to a more interesting dataset, e.g. <https://www.cs.toronto.edu/~kriz/cifar.html>.

by Keith Randall 2013, posted on Stackexchange
n = 50

```
from turtle import *
speed("fastest")
left(90)
forward(3*n)
color("orange", "yellow")
begin_fill()
left(126)
for i in range(5):
    forward(n/5)
    right(144)
    forward(n/5)
    left(72)
end_fill()
right(126)

color("dark.green")
backward(n*4.8)
def thing(d, s):
    if d <= 0: return
    forward(s)
    thing(d-1, s*.8)
    right(120)
    thing(d-3, s*.5)
    right(120)
    thing(d-3, s*.5)
    right(120)
    backward(s)
thing(15, n)
backward(n/2)

import time
time.sleep(60)
```

Happy holidays !